



Special Topics

Semiconductor Fabrication Chemistry: Doping, Photolithography, and Etching in Silicon Chip Manufacturing

Submitted by:

Bangcaya, Marpel Mar B.

Submitted to:

Ms. Nanette D. Santos

CM011L-FOPI01



1. INTRODUCTION

Every modern computing device, from a smartphone processor to a data-center GPU, begins as a single crystal of ultrapure silicon. Turning that plain wafer into a working integrated circuit is fundamentally a chemistry problem: engineers must locally alter the material's electrical, physical, and geometric properties with nanometer-scale precision using controlled chemical reactions rather than mechanical shaping. This report examines the three chemical processes that make this transformation possible: photolithography, which chemically transfers a circuit pattern onto the wafer; doping, which chemically alters silicon's electrical conductivity by introducing impurity atoms; and etching, which chemically removes material to carve out the transistor structures the pattern defines. The objective of this report is to explain the underlying chemistry of these three processes and to connect that chemistry to real fabrication trade-offs — resolution, junction depth, and etch profile — that computer engineers encounter indirectly through chip cost, performance, and availability. The report is limited to these front-end-of-line chemical fabrication steps and does not cover circuit design, chip packaging, or post-fabrication testing.

A. Brief Background

Silicon fabrication traces back to the planar transistor process developed in the late 1950s, when engineers at Fairchild Semiconductor combined oxide masking with selective doping to build transistors on a flat wafer surface instead of shaping each one individually. Since then, the same three chemical operations — photolithography, doping, and etching — have been refined and repeated dozens of times per wafer to build the billions of transistors found in a modern processor. Each new chip generation, commonly called a technology “node,” pushes these chemical processes toward smaller dimensions, shorter light wavelengths, and more precisely controlled reactions. The industry's ability to keep shrinking transistors, an economic and technical trend informally known as Moore's Law, has therefore always depended on chemistry keeping pace with circuit-design ambitions. These same three processes underpin the microprocessors, memory chips, and sensors that computer engineers design around and build systems upon today.

B. Methodology

This report was developed through a structured literature review rather than laboratory experimentation, since the assigned topic concerns an industrial chemical process rather than a bench-scale reaction. Information was gathered from semiconductor process documentation, cleanroom technical guides, and sustainability research published by industry sources such as imec (see References), and each major technical claim was cross-checked across at least two independent sources. The quantitative relationships that describe photolithography resolution, dopant diffusion, and ion-implantation dosing were verified using standard semiconductor-physics formulas and are worked through in the Calculations subsection to show how the underlying chemistry translates into real fabrication design choices.

2. ANALYSIS

A. Calculations

Photolithography, doping, and etching are each governed by well-defined physical relationships. Three representative calculations are worked through below to connect the chemistry of chip fabrication to concrete design numbers.



1. Photolithography resolution (Rayleigh criterion)

The smallest feature a lithography system can print (the critical dimension, CD) follows $CD = k_1 \times (\lambda / NA)$, where λ is the exposure wavelength, NA is the lens's numerical aperture, and k_1 is a process constant (assumed here as 0.3 for both cases so wavelength and NA can be compared directly). Depth of focus follows a related relationship, $DOF = k_2 \times (\lambda / NA^2)$.

- Deep-UV, ArF immersion ($\lambda = 193$ nm, $NA = 1.35$): $CD = 0.3 \times (193 / 1.35) \approx 42.9$ nm

- EUV lithography ($\lambda = 13.5$ nm, $NA = 0.33$): $CD = 0.3 \times (13.5 / 0.33) \approx 12.3$ nm

The roughly 15× shorter EUV wavelength lets it print features about 3.5× smaller than 193 nm immersion lithography even though its numerical aperture is lower, which is precisely why the industry moved to EUV for sub-10 nm-class nodes. (In practice, real EUV depth of focus is further limited by mask 3-D effects and flare, which this simplified model does not capture.)

2. Dopant diffusion depth (Fick's second law, constant-source diffusion)

For a dopant diffusing in from a constant-concentration surface source, the concentration profile is $C(x,t) = C_s \times \text{erfc}[x / (2\sqrt{Dt})]$, where D is the dopant's temperature-dependent diffusion coefficient and t is the diffusion ("drive-in") time. Using a typical order-of-magnitude diffusion coefficient for boron in silicon near 1200°C, $D \approx 2 \times 10^{-13}$ cm²/s, a 30-minute (1800 s) drive-in gives a characteristic diffusion length $\sqrt{Dt} \approx 1.9 \times 10^{-5}$ cm ≈ 0.19 μm — illustrating why furnace diffusion naturally produces junctions tens to hundreds of nanometers deep rather than the few-nanometer junctions modern transistors require.

3. Ion implantation dose

Implant dose is set directly by the ion beam current: $Q = (I \times t) / (q \times A)$, where I is beam current, t is implant time, q is the elementary charge (1.602×10^{-19} C), and A is the implanted area. For $I = 1$ mA delivered for $t = 10$ s over a $A = 100$ cm² wafer: $Q = (0.001 \times 10) / (1.602 \times 10^{-19} \times 100) \approx 6.2 \times 10^{14}$ ions/cm² — within the typical 10^{13} – 10^{15} ions/cm² range used for source/drain and well implants. Because dose and penetration depth are set independently (by beam current/time and accelerating voltage, respectively), ion implantation gives far tighter control than diffusion, which is why it replaced diffusion for the shallow, precisely-doped junctions advanced nodes require.

B. Table

Table 1. Comparison of wet-chemical and dry (plasma) etching in silicon chip fabrication.

Aspect	Wet Etching	Dry (Plasma) Etching
Etchant phase	Liquid chemical solution	Reactive gas plasma / ion beam
Example chemicals	HF, HNO ₃ , H ₃ PO ₄ , HCl; KOH, TMAH (alkaline)	Fluorocarbon, chlorine, and bromine species in plasma (e.g., CF ₄ /SF ₆ -type gases)
Etch profile	Mostly isotropic — rounded sidewalls	Strongly anisotropic — near-vertical sidewalls
Selectivity	High (strong chemical selectivity between materials)	Lower selectivity, traded for geometric precision
Etch rate / equipment	Higher etch rate, simpler equipment	Lower etch rate, complex vacuum/plasma equipment
Typical fab use today	Cleaning, oxide strip, less dimensionally-critical steps	Defining nanometer-scale transistor and interconnect structures



C. Interpretation of Results

The calculations and the comparison table point to the same underlying pattern: as transistors shrink, the industry consistently trades simplicity and cost for chemical and geometric control. The resolution calculation shows that shrinking feature size requires a shorter exposure wavelength, which is why the industry moved from 193 nm immersion lithography to 13.5 nm EUV lithography — at the cost of far more complex, expensive exposure tools and a much shallower depth of focus that forces tighter process control on every other step. The diffusion calculation shows why furnace diffusion, which naturally produces relatively deep, hard-to-control junctions, has been displaced by ion implantation for critical junctions: implantation sets dose and depth independently and precisely, which the dose calculation confirms falls neatly within the 10^{13} – 10^{15} ions/cm² range fabs actually use. The etching comparison shows the same trade-off in a different form — dry plasma etching is slower, less selective, and needs far more complex vacuum equipment than wet etching, but its strong anisotropic control is the only way to carve the near-vertical, nanometer-scale sidewalls that modern transistors and interconnects require, so wet etching has been pushed into cleaning and less dimensionally-critical roles instead.

This same trend carries an environmental cost that is directly traceable to the chemistry involved. Dry etching and chamber cleaning rely on fluorinated gases (including CF₄, SF₆, and NF₃-type chemistries) that have very high global-warming potential, and both lithography and rinse chemistries have historically used PFAS ("forever chemicals") for their resistance to the acids and solvents used elsewhere in the process. Industry data show wet processing alone accounts for roughly half of all chemical use per wafer, and chip production is estimated to generate on the order of 185 million tons of CO₂-equivalent emissions annually (imec, n.d.). This is why leading fabs are investing in abatement and pulsed, lower-consumption etch techniques — one industry pilot cut CF₄ consumption by 98% using pulsed "Transient Assisted Processing" instead of continuous plasma exposure (imec, n.d.) — and in ultrapure-water recycling and PFAS-reduction programs. For a computer engineer, this matters beyond environmental awareness: chemical and regulatory pressure on these processes directly affects fab capacity, chip cost, and where future manufacturing can be sited.

3. CONCLUSION

Building a working silicon chip is, at its core, a sequence of controlled chemical reactions repeated with extraordinary precision: photolithography chemically transfers a circuit pattern using light-sensitive polymers that either break apart or cross-link on exposure; doping chemically rewrites silicon's conductivity through diffusion or ion implantation of boron, phosphorus, or arsenic atoms; and etching chemically carves the resulting structures using either selective wet-chemical solutions or anisotropic dry plasmas. The worked calculations confirmed why the industry has moved the way it has — toward shorter-wavelength EUV lithography, implantation-based doping, and dry plasma etching — as each shift buys tighter dimensional control at the cost of more complex, more chemically demanding, and more environmentally significant processes. This satisfies the report's objective of connecting fabrication chemistry to concrete engineering trade-offs rather than treating chip manufacturing as a black box. For a computer engineering student, the significance is practical as well as academic: the resolution, junction, and etch-profile limits set by this chemistry are the same limits that ultimately determine chip performance, cost, and availability, and the same chemistry's environmental footprint is increasingly a factor in where and how future chips can be produced.



4. References

- Brigham Young University Cleanroom. (n.d.). Photoresist / photolithography process. Retrieved August 29, 2026, from <https://cleanroom.byu.edu/photoresist-photolithography-process>
- Halbleiter.org. (n.d.). Doping techniques. Semiconductor Technology from A to Z. Retrieved August 29, 2026, from <https://www.halbleiter.org/en/waferfabrication/doping/>
- imec. (n.d.). How can we reduce environmental impact in chip manufacturing? Retrieved August 29, 2026, from <https://www.imec-int.com/en/articles/how-can-we-reduce-environmental-impact-chip-manufacturing>
- Wevolver. (n.d.). Positive vs. negative photoresist: A comprehensive guide to photolithography in semiconductor fabrication. Retrieved August 29, 2026, from <https://www.wevolver.com/article/positive-vs-negative-photoresist-a-comprehensive-guide-to-photolithography-in-semiconductor-fabrication>
- Xometry. (n.d.). Dry etching vs. wet etching: Differences and applications. Retrieved August 29, 2026, from <https://www.xometry.com/resources/blog/dry-etching-vs-wet-etching/>